

# Neural Network Exercise


Kim Mingi (xp2425@khu.ac.kr)


## Reference



A. Wolfram Tutorial “Introduction to Neural Nets” <https://reference.wolfram.com/language/tutorial/NeuralNetworksIntroduction.html#1969481202>



- MNIST handwriting examples

```
In[1]:= ResourceSearch["MNIST"]
```

```
Out[1]= {ResourceObject[ Name: FashionMNIST »  
Type: DataResource  
Description: A small MNIST-like fashion product image dataset
```

```
ResourceObject[ Name: MNIST »  
Type: DataResource  
Description: Database of handwritten digits commonly used for training image processes
```

```
ResourceObject[ Name: LeNet Trained on MNIST Data »  
Type: NeuralNet  
Description: Identify the handwritten digit in an image
```

```
ResourceObject[ Name: CapsNet Trained on MNIST Data »  
Type: NeuralNet  
Description: Identify the handwritten digit in an image
```

```
In[2]:= trainingData = ResourceData["MNIST", "TrainingData"]
```

Out[2]=

```
{
  0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0,
  0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0,
  0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0,
  0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0,
  0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0, 0 → 0,
  0 → 0, 0 → 0, ... 59910 ..., 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9,
  9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9,
  9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9,
  9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9,
  9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9, 9 → 9}

```

large output

[show less](#)

[show more](#)

[show all](#)

[set size limit...](#)

```
In[3]:= testData = ResourceData["MNIST", "TestData"]
```

Out[3]=

large output | [show less](#) | [show more](#) | [show all](#) | [set size limit...](#)

Pick a few random examples

```
In[4]:= RandomSample[trainingData, 3]
```

Out[4]= { 0 → 0, 2 → 2, 3 → 3 }

## 1. Introduction to NN layers

Layers in Mathematica

In[15]:= ? \*Layer


## ▼ System`

AggregationLayer	LongShortTermMemoryLayer
BasicRecurrentLayer	MeanAbsoluteLossLayer
BatchNormalizationLayer	MeanSquaredLossLayer
CatenateLayer	PaddingLayer
ConstantArrayLayer	PartLayer
ConstantPlusLayer	PoolingLayer
ConstantTimesLayer	ReplicateLayer
ContrastiveLossLayer	ReshapeLayer
ConvolutionLayer	ResizeLayer
CrossEntropyLossLayer	SequenceAttentionLayer
DeconvolutionLayer	SequenceLastLayer
DotLayer	SequenceMostLayer
DotPlusLayer	SequenceRestLayer
DropoutLayer	SequenceReverseLayer
ElementwiseLayer	SoftmaxLayer
EmbeddingLayer	SpatialTransformationLayer
FlattenLayer	SummationLayer
GatedRecurrentLayer	ThreadingLayer
ImageAugmentationLayer	TotalLayer
InstanceNormalizationLayer	TransposeLayer
LinearLayer	UnitVectorLayer
LocalResponseNormalizationLayer	

Create a LinearLayer that takes as input a vector of length 2 and produces as output a vector of length 3

In[18]:= **linear** = LinearLayer[3, "Input" → 2]

Out[18]= LinearLayer[



**Parameters**  
OutputDimensions: 3

---

**Arrays**  
Weights: matrix (size: 3 × 2)  
Biases: optional vector (size: 3)

---



**Ports**  
Input: vector (size: 2)  
Output: vector (size: 3)

In[19]:= **linear**[{2, 3}]

Out[19]= \$Failed

Above layer is an ‘uninitialized layer’ which is to say, the learning is not yet started. By using NetInitialize, we can initialize the layer

```
In[20]:= initLinear = NetInitialize[linear]
```

```
Out[20]= LinearLayer[  Parameters  
OutputDimensions: 3  
Arrays  
Weights: matrix (size: 3 × 2)  
Biases: vector (size: 3)  
Ports  
Input: vector (size: 2)  
Output: vector (size: 3)
```

```
In[22]:= initLinear[{{2, 3}, {2, 4}, {1, 2}}]
```

```
Out[22]= {{2.38762, 1.40096, 0.639938}, {3.027, 1.76667, 0.264892}, {1.5135, 0.883337, 0.132446}}
```

Without pesky calculations, we can easily 'extract' the weight and biases.

```
In[27]:= W = NetExtract[initLinear, "Weights"];  
W // MatrixForm
```

```
Out[28]//MatrixForm= 
$$\begin{pmatrix} 0.234746 & 0.639378 \\ 0.151909 & 0.365714 \\ 0.882538 & -0.375046 \end{pmatrix}$$

```

```
In[24]:= NetExtract[initLinear, "Biases"]
```



```
Out[24]= {0., 0., 0.}
```

```
In[29]:= Dot[W, {2, 3}]
```

```
Out[29]= {2.38762, 1.40096, 0.639938}
```

LinearLayer has only one input, on the other hand, some layers have more than two inputs; MeanSquaredLossLayer needs the *input* and the *target* as two inputs

```
In[30]:= msloss = MeanSquaredLossLayer[]
```

```
Out[30]= MeanSquaredLossLayer[  Parameters  
none  
Ports  
Input: tensor  
Target: tensor  
Loss: real
```

Apply the layer to two inputs

```
In[31]:= msloss[<|"Input" → {1, 2, 3}, "Target" → {4, 0, 4}|>]
```

```
Out[31]= 4.66667
```

## 2. Net Encoders & Net Decoders

Before we train or test the data, it should be transformed to numeric tensor, which is a NetEncoder's job; to translate the data to numeric tensors. Let's see some of the simple examples




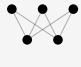
```
In[41]:= imageenc[img] // Dimensions
```

```
Out[41]= {1, 12, 12}
```

We can attach encoder to a layer

```
In[42]:= pool = PoolingLayer[10, "Input" → NetEncoder[{"Image", 64}]]
```

```
Out[42]= PoolingLayer[
```

		<b>Parameters</b>
		KernelSize: {10, 10}
		Stride: {1, 1}
		PaddingSize: {0, 0}
		Function: Max
		Dimensionality: 2
		<b>Ports</b>
		Input: image
		Output: 3-tensor (size: 3 × 55 × 55)

```
]
```

Then we can directly apply the image the layer

```
In[48]:= pool[img] // Shallow
```

```
Out[48]/Shallow= {{ { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> },
  { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, <<45>> },
  { { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> },
  { <<55>> }, { <<55>> }, { <<55>> }, <<45>> }, { { <<55>> }, { <<55>> }, { <<55>> },
  { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, { <<55>> }, <<45>> } }
```

Convert the output back to an image

```
In[51]:= Image[%, Interleaving → False]
```

```
Out[51]= 
```



Finally create the final image encoder for MNIST, whose grayscale images are of 28×28.

```
In[64]:= mnistEnc = NetEncoder[{"Image", {28, 28}, "ColorSpace" → "Grayscale"}]
```

```
Out[64]= NetEncoder[
```

Type:	Image
ImageSize:	{28, 28}
ColorSpace:	Grayscale
ColorChannels:	1
MeanImage:	None
VarianceImage:	None
Output:	3-tensor (size: 1 × 28 × 28)

```
]
```

```
In[54]:= trainingData[[1, 1]] // ImageDimensions
```

```
Out[54]= {28, 28}
```

Next, create a decoder for MNIST handwritings. Conversely, decoder translate the numeric tensor to some class. In case of MNIST, there are ten classes, from 0 to 9.

```
In[56]:= mnistDec = NetDecoder[{"Class", Range[0, 9]}]
```

```
Out[56]= NetDecoder[
  Type:      Class
  Labels:    {<<1>>}
  Dimensions: 10
]
```

```
In[58]:= mnistDec[{0, 0, 0, 0, 0, 0, 0, 0, 0, 1}]
```

```
Out[58]= 9
```

```
In[60]:= mnistDec[{0.1, 0.3, 0.6, 0, 0, 0, 0, 0, 0, 0}, "TopProbabilities"]
```


```
Out[60]= {2 → 0.6, 1 → 0.3, 0 → 0.1}
```

```
In[61]:= mnistDec[{0.1, 0.3, 0.6, 0, 0, 0, 0, 0, 0, 0}, "Entropy"]
```

```
Out[61]= 0.897946
```

Similarly to encoder, we can attach a decoder to the layer

```
In[62]:= pool = PoolingLayer[10, "Input" → NetEncoder[{"Image", 64}], "Output" → NetDecoder["Image"]]
```

```
Out[62]= PoolingLayer[
  + 
  Parameters
  KernelSize: {10, 10}
  Stride: {1, 1}
  PaddingSize: {0, 0}
  Function: Max
  Dimensionality: 2
]
```

```
In[63]:= pool[img]
```

```
Out[63]= 
```

Now construct NN to concatenate the layers together.



```
In[65]:= uninitializedLenet =
  NetChain[{ConvolutionLayer[20, 5], ElementwiseLayer[Ramp], PoolingLayer[2, 2],
    ConvolutionLayer[50, 5], ElementwiseLayer[Ramp], PoolingLayer[2, 2],
    FlattenLayer[], LinearLayer[500], ElementwiseLayer[Ramp], LinearLayer[10],
    SoftmaxLayer[]}, "Input" → mnistEnc, "Output" → mnistDec]
```

Out[65]= NetChain [

		image
Input		3-tensor (size: $1 \times 28 \times 28$ )
1	ConvolutionLayer	3-tensor (size: $20 \times 24 \times 24$ )
2	Ramp	3-tensor (size: $20 \times 24 \times 24$ )
3	PoolingLayer	3-tensor (size: $20 \times 12 \times 12$ )
4	ConvolutionLayer	3-tensor (size: $50 \times 8 \times 8$ )
5	Ramp	3-tensor (size: $50 \times 8 \times 8$ )
6	PoolingLayer	3-tensor (size: $50 \times 4 \times 4$ )
7	FlattenLayer	vector (size: 800)
8	LinearLayer	vector (size: 500)
9	Ramp	vector (size: 500)
10	LinearLayer	vector (size: 10)
11	SoftmaxLayer	vector (size: 10)
Output		class

(uninitialized)

**ElementwiseLayer (2)**

**Parameters**

Function: Ramp

---


**Ports**

Input: 3-tensor (size:  $20 \times 24 \times 24$ )

Output: 3-tensor (size:  $20 \times 24 \times 24$ )

Note: Ramp: ReLU

```
In[84]:= Plot[Ramp[x], {x, -5, 10}]
```



```
In[66]:= lenet = NetInitialize[uninitializedLenet]
```

Out[66]= NetChain [

		image
Input		3-tensor (size: $1 \times 28 \times 28$ )
1	ConvolutionLayer	3-tensor (size: $20 \times 24 \times 24$ )
2	Ramp	3-tensor (size: $20 \times 24 \times 24$ )
3	PoolingLayer	3-tensor (size: $20 \times 12 \times 12$ )
4	ConvolutionLayer	3-tensor (size: $50 \times 8 \times 8$ )
5	Ramp	3-tensor (size: $50 \times 8 \times 8$ )
6	PoolingLayer	3-tensor (size: $50 \times 4 \times 4$ )
7	FlattenLayer	vector (size: 800)
8	LinearLayer	vector (size: 500)
9	Ramp	vector (size: 500)
10	LinearLayer	vector (size: 10)
11	SoftmaxLayer	vector (size: 10)
Output		class

```
In[78]:= samples = RandomSample[trainingData, 10]
```

```
Out[78]= { 0 → 0, 6 → 6, 1 → 1, 1 → 1,
           1 → 1, 4 → 4, 4 → 4, 0 → 0, 1 → 1, 9 → 9 }
```

```
In[81]:= samplelist = Table[samples[[i, 1]], {i, 10}]
```

```
Out[81]= { 0, 6, 1, 1, 1, 4, 4, 0, 1, 9 }
```

```
In[82]:= lenet[samplelist]
```

```
Out[82]= {8, 8, 8, 7, 8, 8, 8, 8, 8, 8}
```

Pretty stupid...